

# Supervising Your AI With The Red Team

We have supervised machine learning as a concept but supervising artificial intelligence (AI) itself has not been deeply thought out. In this article we focus on ongoing controls over the AI solution and see what involves supervising AI and elaborate on the concept of a red team in the context of AI solutions



**Anand Tamboli**  
is a serial entrepreneur, speaker, award-winning published author and emerging technology thought leader

**I**n military parlance, the red team is one that uses its skills to imitate an attack and the techniques a potential enemy might use. The other group, which acts as the defence, is called the blue team. Cybersecurity has picked up these terms for their use, where they signify the same functionality.

A red team is an independent group of people that challenges an organisation to improve its defenses and effectiveness by acting as an adversary or attacker. In layman's terms, it is a form of ethical hacking—a way to test how well an organisation would do in the face of a real attack.

If effectively implemented, red teams can expose vulnerabilities and risks not only in the technology infrastructure but also in people and physical assets.

You must have heard the saying “The best offence is a good defense.” Having a red team is the right step in putting up this good defense. Whether it is a complex artificial intelligence (AI) solution or merely some basic technology solution, red-teaming can give a competitive edge to the business.

A comprehensive red team exercise may involve penetration testing (also known as pen testing), social engineering and physical

intrusion. Either all of these could be carried out, or a combination of those in the right manner as the red team sees fit to expose vulnerabilities.

If there exists a red team, then there also must exist a blue team. This assumption stems from the premise that system developments are done inhouse only. However, with AI systems, this can change, and the actual blue team may be the technology vendor.

## Building the red team

Ideally, the red team needs at least two people to be effective, though many teams have up to five. In case of a large company, fifteen-plus members might be needed in the team to work on several fronts.

Depending on the type of AI that is deployed, the red team's composition might change with different skillsets—structure is necessary to maximise the team's effectiveness.

Typically, you would need physical security experts, such as the ones who understand and can deal with physical locks, door codes and other aspects. You would also need a social engineer to phish out information through emails, phone calls,